

When Graph Structure Becomes a Liability: A Critical Re-Evaluation of Graph Neural Networks for Bitcoin Fraud Detection under Temporal Distribution Shift

Saket Maganti

Abstract

Published results on the Elliptic Bitcoin Dataset place GCN, GraphSAGE, GAT, and EvolveGCN above feature-only baselines, and we show this ranking is a protocol artefact rather than a property of graph learning. Under a strictly inductive protocol (encoder trained on the time-step ≤ 34 relabeled subgraph, evaluated on the full graph at inference) with 10 seeds and per-timestep reporting, Random Forest on raw 165-dimensional features reaches $F1 = 0.821$ and beats every GNN we tested; GraphSAGE, the strongest graph encoder, reaches only $F1 = 0.689 \pm 0.017$. A paired controlled experiment holding architecture, optimiser, loss, and seed constant across 10 matched seeds shows GraphSAGE scoring $F1 = 0.294 \pm 0.028$ when trained transductively and $F1 = 0.689 \pm 0.017$ when trained inductively—a 39.5-point paired gap (Cohen’s $d = 15.8$, $p = 2.6 \times 10^{-12}$) explained entirely by training-time exposure to test-period adjacency. Earlier drafts reported $F1 = 0.807$ for a concatenation hybrid of GraphSAGE embeddings and raw features; under the clean protocol the same hybrid falls to $F1 = 0.699 \pm 0.015$, and the GNN contributes a statistically reliable but small $+0.018$ F1 lift over a matched-capacity MLP substitute ($p = 0.015$, $d = +1.20$) that is dwarfed by the 0.124 F1 gap to raw features alone. A 10-seed edge-shuffle ablation further shows that on Elliptic *randomly shuffled edges outperform the real transaction graph* by 8.9 F1 points, so the dataset’s topology is not merely uninformative but actively harmful under its sparse, prior-shifted conditions. Code, checkpoints, 10-seed per-run artefacts, and the strict-inductive protocol are released so future work can be compared against a leakage-free baseline.

1 Introduction

Graph Neural Networks (GNNs) are the default architecture for fraud detection on transaction graphs, and the motivation is sound. Organised fraud rarely happens through isolated transactions: laundering requires layering across wallets, fraud rings share infrastructure, and mixing services leave topologically distinctive flow patterns [Akoglu et al., 2015]. The Elliptic Bitcoin Dataset [Weber et al., 2019], released in 2019, labels 46,564 Bitcoin transactions across 49 temporal steps and has become the canonical benchmark for this hypothesis. Published results line up neatly behind it: GCN ($F1 = 0.70$ [Weber et al., 2019]), augmented GCN ($F1 = 0.74$ [Alarab et al., 2020]), GraphSAGE with self-supervised pre-training ($F1 = 0.75$ [Lo et al., 2023]), and EvolveGCN ($F1 = 0.77$ [Pareja et al., 2020]) all beat the feature-only baselines reported alongside them.

We show that this ranking is an artefact of the evaluation protocol. Every prior Elliptic study trains with *transductive* message passing, in which the full graph—including test-period nodes and their edges—is visible at every training forward pass. Several “inductive” setups in the same literature run the encoder on the full graph and simply mask the loss, which still leaks test-period feature statistics through batch normalisation and neighbourhood aggregation. Every prior study also reports a single aggregate F1 over the entire 15-step test window, concealing a near-total collapse after step 42 when the fraud base rate drops $39\times$ from

11.6% to 0.3%. This paper tightens both choices. We train on the relabeled subgraph induced by time-step ≤ 34 , so neither message passing nor batch statistics can observe any test-period vector during training; we report per-timestep F1 alongside aggregates; and we run 10 seeds with 95% bootstrap confidence intervals for every headline number. Figure 2 in Section 7.1 shows the per-timestep F1 picture that the aggregate metric conceals.

The protocol changes are minor; the empirical consequences are not. Three findings drive the paper.

Under strict inductive protocol the ranking inverts. Random Forest on the 165 raw features reaches $F1 = 0.821 \pm 0.003$ (10 seeds), above every GNN we tested. GraphSAGE, the strongest graph encoder under strict inductive, reaches $F1 = 0.689 \pm 0.017$ (10 seeds) and in turn beats the 3-layer MLP at $F1 = 0.549 \pm 0.015$ (10 seeds). This *reverses* the MLP>GNN ranking produced by random-partition mini-batching, which leaks test-period features through each batch’s induced subgraph. A paired controlled experiment quantifies the leakage directly: across 10 matched seeds GraphSAGE scores $F1 = 0.294 \pm 0.028$ when trained transductively and $F1 = 0.689 \pm 0.017$ when trained on the inductive subgraph—a 39.5-point paired gap (paired $t = 49.97$, $p = 2.6 \times 10^{-12}$, Cohen’s $d = 15.8$), with the inductive model *stronger* than its transductive counterpart (Section 7.7).

The hybrid result does not survive the clean protocol. Earlier drafts reported that a concatenation hybrid of GraphSAGE embeddings and raw features reaches $F1 = 0.807$ through a downstream Random Forest. Under strict inductive training the same hybrid drops to $F1 = 0.699 \pm 0.015$ over 10 seeds. A paired ablation that swaps the GNN for an MLP of identical capacity shows graph structure contributes a statistically reliable but small lift ($+0.018$ F1, Welch $p = 0.015$, $d = +1.20$), dwarfed by the 0.124 F1 gap to raw features alone ($F1 = 0.823 \pm 0.002$, Section 7.8). Whatever lift the Random Forest recovers from the 256-dimensional embedding on top of raw features is real but practically inconsequential; raw features remain the strongest input representation.

On Elliptic, random edges beat real edges. A multi-seed graph-structure ablation (10 seeds) shows that randomly shuffling the transaction edges *improves* GraphSAGE from $F1 = 0.290 \pm 0.019$ to $F1 = 0.380 \pm 0.028$ ($+0.089$ F1), and that removing the edges entirely still beats the real graph (0.316 ± 0.018 , $+0.025$ F1). Under this dataset’s sparse, prior-shifted conditions, the real topology is not merely uninformative: its systematically mixed neighbourhoods (each fraud node surrounded by licit counterparties) actively pull fraud representations toward the licit manifold under shift.

Contributions. (i) A strictly inductive, 10-seed benchmark for four GNN architectures plus MLP and Random Forest baselines on Elliptic, with 95% bootstrap CIs, per-timestep F1, and full release of code and checkpoints (Section 7.1). (ii) A direct transductive-vs.-inductive comparison holding architecture, optimiser, loss, and seed constant across 10 matched seeds, quantifying a 39.5-point paired F1 gap ($p = 2.6 \times 10^{-12}$, $d = 15.8$) explained entirely by training-time exposure to test-period adjacency (Section 7.7). (iii) A calibrated re-assessment of the widely reported “GNN embeddings plus raw features” hybrid recipe: under strict inductive the same hybrid scores 0.699 ± 0.015 and graph embeddings add a small, statistically reliable lift over a matched-capacity MLP substitute ($+0.018$ F1, $p = 0.015$, $d = +1.20$) that is dwarfed by the 0.124 F1 gap to raw features alone, demonstrating that the 0.807 number from earlier drafts was a protocol artefact and that raw features remain the strongest input representation (Section 7.8). (iv) An edge-shuffle ablation showing that on Elliptic, random wiring beats the real transaction graph, together with a mechanistic explanation grounded in neighbourhood composition under prior shift (Section 7.6). (v) A prior-work audit of the most-cited Elliptic GNN results and the classical baselines reported alongside them, identifying exactly which protocol shortcut each relies on, so future work can be benchmarked against a leakage-free baseline (Section 2.1).

2 Related work

We organise this section around a single observation: to our knowledge, every prior Elliptic study we surveyed shares two methodological assumptions whose interaction with the dataset’s temporal structure warrants closer examination. The assumptions are (i) that a transductive protocol, in which the full graph is visible at training time, is an acceptable proxy for deployment, and (ii) that a single aggregate F1 over the 15-step test window adequately summarises temporal behaviour. Our results suggest that neither assumption holds well under the distribution shift present in Elliptic.

2.1 GNN-based fraud detection on Elliptic

Graph-based anomaly detection predates deep learning. Early approaches used belief propagation and Markov random fields to identify fraudulent accounts in payment graphs [Akoglu et al., 2015]. Differentiable message-passing architectures [Gilmer et al., 2017] made it possible to learn the aggregation function end-to-end, and every subsequent Elliptic study inherits this framing without revisiting whether end-to-end graph learning is appropriate for the specific structure of the Bitcoin transaction graph.

Table 1 audits the four most-cited prior Elliptic GNN studies—Weber, Alarab, Pareja, and Lo—together with the two classical baselines (Logistic Regression, Random Forest) reported alongside them, along three axes: (a) does the encoder’s message passing observe test-period edges or features during training, (b) does the reported metric aggregate across all 15 test steps or break out per-step behaviour, and (c) how many random seeds back the headline F1. Every GNN result we audit trains with full-graph message passing (four of four) or, in setups labelled “inductive,” masks the loss while still running the encoder forward pass on the full graph—a protocol that still leaks batch statistics through the aggregator. All six reported numbers aggregate over the full test window, and five of the six are single-seed or single-run. This does not mean any individual result is invalid; it means the reported numbers are not directly comparable to ours under strict inductive protocol with 10-seed variance, and the gap we document should be read as a protocol-induced gap, not an architecture-induced gap.

Table 1: Prior-work audit: Elliptic GNN results and the evaluation shortcuts each relies on. “Train msg passing” asks whether the training-time forward pass observes test-period edges/features in aggregation or batch normalisation; “Reported” is the headline F1 as published; “Per-step” indicates whether per-timestep F1 is reported. Our inductive SAGE number is directly comparable to none of these, which is the point: the protocol gap explains most of the apparent architecture gap.

Study	Model	Train msg passing	Reported F1	Per-step?	Seeds
Weber et al. [Weber et al., 2019]	GCN	Transductive	0.70	No	1
Alarab et al. [Alarab et al., 2020]	GCN+feat	Transductive	0.74	No	1
Pareja et al. [Pareja et al., 2020]	EvolveGCN	Snapshots (trans.)	0.77	No	1
Lo et al. [Lo et al., 2023]	GraphSAGE + SSL	Transductive	0.75	No	3
Weber et al. [Weber et al., 2019]	RF (raw features)	N/A	0.64	No	1
Weber et al. [Weber et al., 2019]	LR (raw features)	N/A	0.53	No	1
Ours	GraphSAGE (strict ind.)	Inductive	0.688 ± 0.016	Yes	10
Ours	MLP (strict ind.)	N/A	0.549 ± 0.015	Yes	10
Ours	RF (raw features)	N/A	0.821 ± 0.003	Yes	10

Weber et al. [Weber et al., 2019] established the founding GCN baseline ($F1 = 0.70$) and reported it alongside Logistic Regression (0.53) and Random Forest (0.64) under a transductive protocol. Their own paper notes that the 71 manually engineered aggregate features already provide a surrogate for one round of message passing, which bounds the additional gain available to a learned GCN over a tuned MLP; subsequent work cites the 0.70 number without quoting the caveat. Pareja et al. [Pareja et al., 2020] introduced EvolveGCN, which updates GCN weight

matrices via a GRU across time snapshots, and reported $F1 = 0.77$ on Elliptic; the result requires full GPU training over every one of the 34 training snapshots, and—critically—the evaluation aggregates over all 15 test steps without reporting per-step behaviour, so the collapse after step 42 is invisible in their headline metric. Alarab et al. [Alarab et al., 2020] combined GCN with additional feature engineering and class-weighted loss for $F1 = 0.74$, again under transductive evaluation. Lo et al. [Lo et al., 2023] applied GraphSAGE with structural augmentation and self-supervised pre-training, reaching $F1 = 0.75$, with the same protocol.

A recurring pattern emerges: each paper adds a technique (temporal weight evolution, pre-training, class weighting) to the transductive baseline and reports a marginal gain, without revisiting the evaluation protocol itself. Our paired controlled experiment (Section 7.7) quantifies what this protocol choice is worth: training GraphSAGE transductively versus inductively, with everything else held constant, changes $F1$ by 39.5 points (0.294 ± 0.028 vs. 0.689 ± 0.017 across 10 matched seeds, paired $p = 2.6 \times 10^{-12}$), so the protocol change alone accounts for considerably more than the architecture choice. Under strict inductive evaluation, EvolveGCN’s 0.77 and Lo et al.’s 0.75 cannot be compared directly to our strict-inductive SAGE number because the published numbers include transductive leakage; the appropriate apples-to-apples comparison would re-run each published method under our protocol, which we leave for future work.

Outside Elliptic, GNN fraud detection has been applied to e-commerce at scale [Wang et al., 2019], financial networks [Liu et al., 2019], and insurance claims [Pourhabibi et al., 2020]. The shared assumption across these studies is distributional stability between training and deployment. That assumption is rarely stated and almost never tested; in Elliptic it breaks down sharply, and there is little reason to expect it to hold in the adversarial settings for which these systems are typically deployed.

2.2 Classical machine learning for fraud

Before GNNs, classical methods set the standard for transaction fraud detection [Ngai et al., 2011, Hilal et al., 2022]. Weber et al. [Weber et al., 2019] reported Logistic Regression ($F1 = 0.53$) and Random Forest ($F1 = 0.64$) as baselines under transductive evaluation. Gradient-boosted trees [Chen and Guestrin, 2016] perform in a similar range on comparable imbalanced tasks. These classical results matter for context: as we show, a Random Forest on the raw features alone, evaluated under strict-inductive protocol, beats every published GNN result on this dataset despite their use of a more permissive (transductive) evaluation.

2.3 Class imbalance in graph learning

Fraud datasets are severely imbalanced: legitimate transactions outnumber fraudulent ones by 7–10 $\times$ in Elliptic and often much more in production. Standard cross-entropy allocates most gradient to the majority class, leaving little incentive to detect rare fraud.

Chawla et al. [Chawla et al., 2002] proposed SMOTE, which synthesises minority samples by interpolating between existing examples in feature space. Zhao et al. [Zhao et al., 2021] adapted this to graphs by interpolating in embedding space and rewiring edges for synthetic nodes. Our approach differs: we clone structurally complete 1-hop ego-graphs around fraud seed nodes, preserving the dense intra-fraud connectivity that GNNs should learn to exploit.

Focal loss [Lin et al., 2017] down-weights easy examples through a modulating factor $(1 - p_t)^\gamma$, concentrating gradient on borderline cases. Square-root class weighting, which we use, provides similar benefits with lower variance than linear weighting.

2.4 Temporal graph learning and distribution shift

EvolveGCN [Pareja et al., 2020] is the standard temporal GNN baseline for Elliptic. DySAT [Sankar et al., 2020] and TGAT [Xu et al., 2020] represent more recent temporal

attention architectures. ROLAND [You et al., 2022] provides a framework for incremental graph learning with live updates. All require GPU hardware and access to full temporal sequences.

Concept drift, the statistical phenomenon where the target distribution changes over time [Gama et al., 2014, Lu et al., 2019], is endemic to fraud detection. Fraudsters adapt in response to detection; prices shift; law enforcement disrupts specific criminal ecosystems. The Elliptic test period captures exactly this kind of abrupt, externally driven drift [Quiñonero-Candela et al., 2009].

2.5 Calibration

Guo et al. [Guo et al., 2017] showed that modern neural networks are systematically overconfident and proposed temperature scaling as a post-hoc fix: divide logits by a scalar $T > 1$ before the softmax. Temperature scaling is rank-preserving: it changes probability values but not their ordering, so the argmax prediction at a fixed threshold is unchanged [Platt, 1999, Zadrozny and Elkan, 2001]. For our analysis, this means temperature scaling can only improve F1 by moving miscalibrated probabilities across the decision threshold. We show this does not happen during temporal collapse.

3 Background

3.1 Graph neural networks

A GNN operates on a graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ with node feature matrix $\mathbf{X} \in \mathbb{R}^{|\mathcal{V}| \times d}$. Each layer updates node representations by aggregating information from neighbours. In the general message-passing framework [Gilmer et al., 2017]:

$$\mathbf{m}_v^{(l)} = \bigoplus_{u \in \mathcal{N}(v)} \phi(\mathbf{h}_v^{(l)}, \mathbf{h}_u^{(l)}, \mathbf{e}_{vu}), \quad \mathbf{h}_v^{(l+1)} = \psi(\mathbf{h}_v^{(l)}, \mathbf{m}_v^{(l)}) \quad (1)$$

where $\bigoplus$ is a permutation-invariant aggregation (mean, sum, max), ϕ is the message function, ψ is the update function, and $\mathcal{N}(v)$ denotes the neighbours of node v . GCN, GraphSAGE, and GAT each instantiate this framework with different choices for ϕ , $\bigoplus$, and ψ .

3.2 Distribution shift and its geometric consequence

Let $P_{\text{train}}(X, Y)$ denote the joint distribution at training time and $P_{\text{test}}(X, Y)$ at deployment. Distribution shift occurs when $P_{\text{train}} \neq P_{\text{test}}$ [Ben-David et al., 2010, Sugiyama et al., 2007]. For fraud detection on Elliptic, the shift is overwhelmingly in the prior: $P_{\text{test}}(Y=\text{fraud})$ drops by more than an order of magnitude while the class-conditional density $P(X | Y)$ is approximately preserved—a textbook case of prior probability shift [Quiñonero-Candela et al., 2009]. We verify this directly in Section 7: MMD between each test step and training is nearly flat (0.20–0.31) even through the collapse region, so the feature distribution itself is stable; what moves is the label mass.

The geometric consequence makes this form of shift operationally severe. A discriminative classifier trained under the 11.6% training prior fits a decision boundary Bayes-optimal for that prior. Deployed at a 0.3% fraud rate, the Bayes-optimal boundary would need to move sharply toward the fraud class, but the learned boundary does not move—its weights were frozen at the end of training. Every borderline fraud node under the old prior is now on the wrong side of a boundary calibrated to an environment that no longer exists.

Three properties follow. First, the collapse is asymmetric: recall drops while precision on existing positives holds up. Second, post-hoc temperature scaling [Guo et al., 2017] cannot fix it: temperature scaling is monotone, preserving the rank of predictions at any fixed threshold, and the problem is precisely that the correct threshold has moved. Third, when the number of fraud

nodes per step shrinks to single digits (step 46: three nodes out of 960), any threshold-based correction is statistically uninformative.

3.3 Transductive vs. inductive evaluation

Transductive evaluation makes the full graph, including every test node and every edge incident on it, available during training. A GNN trained this way aggregates over test-period neighbourhoods during forward passes, and gradients flow through paths that connect unlabelled test nodes to labelled training nodes. Because Elliptic’s 49 time steps form disjoint connected components, the test-period neighbourhood of a test node consists almost entirely of *other test-period nodes*; a transductive model therefore receives messages that aggregate features from nodes whose labels would not exist in a live fraud pipeline. Under severe temporal shift, this exposure can become a significant confound in the evaluation signal.

Inductive evaluation withholds all test-period structure from training. At inference the model processes nodes whose neighbourhood context it has never observed—the setting any deployed fraud detector actually faces. Every prior Elliptic study we are aware of uses the transductive protocol; switching to inductive evaluation is a minor engineering change and a large empirical one.

4 Dataset and problem setup

4.1 The Elliptic Bitcoin Dataset

Elliptic, a blockchain analytics company, assembled this dataset from the Bitcoin transaction ledger and released it publicly in 2019 [Weber et al., 2019]. Each node is a single Bitcoin transaction. Directed edges run from the outputs of one transaction to the inputs of another, encoding BTC flow; following prior work we treat them as undirected for bidirectional message passing. Table 2 gives the summary statistics.

Each of the 165 node features describes directly observable transaction properties. Features 0–93 are *local*: input/output counts, transaction amounts, fees, timing, and statistical aggregates. Features 94–164 are *aggregate*: summaries of local features of immediate neighbours, computed by the dataset creators as a manual proxy for one round of GNN message passing. This design means the MLP already has access to first-order neighbourhood information through fixed aggregate features; the GNN’s additional contribution is learned, dynamic propagation over the graph topology.

Table 2: Elliptic Bitcoin Dataset summary. The 7.6:1 training imbalance and the variation in test-period fraud rates are the two properties that most shape our results.

Property	Value
Nodes (transactions)	203,769
Edges (BTC flows, undirected)	234,355
Node features	165 (94 local + 71 aggregate)
Labeled nodes	46,564 (22.8 %)
Illicit	4,545 (9.8 % of labeled)
Licit	42,019 (90.2 % of labeled)
Unknown nodes (excluded)	157,205
Temporal steps	49 ($\approx$ 2 weeks each)
Training (steps 1–34)	29,894 labeled; 11.6 % illicit
Test (steps 35–49)	16,670 labeled; 6.5 % overall
Training imbalance ratio	7.6:1 (licit:illicit)

Figure 1 makes two things visible that the summary table cannot convey: the temporal split separates a comparatively steady training regime from a volatile test regime, and the most useful features are related but not interchangeable.




Figure 1: Dataset characterisation. *Top left:* illicit rate (train/test split at step 34/35). *Top right:* node counts. *Bottom left:* class composition. *Bottom right:* top-10 feature correlations.

4.2 Why the test period is problematic

Table 3 quantifies the fraud rate variation across test time steps. Steps 35, 38, 41, and 42 carry rates between 10 % and 15 %, close to the training distribution of 11.6 %. Step 46 has just three illicit nodes among 960 labeled transactions (0.3 %). Steps 43–49 average 3.3 % illicit, a $3.5\times$ reduction from training.

This variation reflects genuine changes in Bitcoin fraud activity during early 2019. Major exchanges tightened anti-money-laundering controls, law enforcement closed several high-profile darknet markets, and the 2018–2019 bear market reduced incentives for certain fraud types. The result is a fraud-rate trajectory that drops by $39\times$ between training and the worst test steps.

4.3 Graph construction variants

The Elliptic transaction graph is sparse: each transaction has 2.3 neighbours on average (Table 4), and the 49 connected components correspond to the 49 temporal slices with no edges crossing slice boundaries. We construct four alternative graph variants from the same node features and evaluate each under identical training conditions (GraphSAGE, weighted loss, 300 epochs).

Similarity graph. Cosine similarity between all node pairs per timestep; edge where similarity exceeds 0.92. Yields 3.6M edges with mean degree 17.7.

Feature k-NN graph. Each node connects to its five nearest neighbours in feature space (L2 distance) within the same timestep: 1.0M edges, mean degree 4.9.

Temporal graph. For adjacent timesteps t and $t+1$, each node in t links to its three nearest feature-space neighbours in $t+1$. This encodes temporal continuity across consecutive slices: 1.2M edges, mean degree 5.9.

Augmented graph. Original transaction edges combined with all similarity edges: 4.1M edges, mean degree 19.9.

Table 3: Illicit rate and node counts per test time step. Steps 35–42 maintain rates broadly comparable to training. Steps 43–49 drop to near-zero fraud rates.

Step	Illicit %	N	N illicit	N licit	Region
35	13.6	1,088	148	940	Stable
36	1.9	660	13	647	Low
37	8.0	750	60	690	Moderate
38	14.7	1,224	180	1,044	Stable (peak)
39	6.8	882	60	822	Moderate
40	9.2	1,087	100	987	Moderate
41	10.2	1,138	116	1,022	Stable
42	11.1	1,045	116	929	Last stable step
43	1.8	997	18	979	Collapse onset
44	1.5	960	14	946	Collapsed
45	0.4	940	4	936	Near-zero
46	0.3	960	3	957	Minimum
47	2.6	980	25	955	Collapsed
48	7.6	986	75	911	Partial recovery
49	11.8	953	112	841	Partial recovery
Training avg	11.6		3,462	26,432	—

Table 4: Topology of each graph variant. Clustering estimated from a 3,000-node sample.

Graph type	Edges	Mean deg.	Max deg.	Clustering	CC
Original	468,710	2.30	473	0.000	49
Similarity	3,611,346	17.72	482	0.045	—
Feature k-NN	1,007,816	4.95	25	0.000	—
Temporal	1,207,890	5.93	331	0.000	—
Augmented	4,057,870	19.91	484	0.045	—

The original graph has clustering coefficient zero: no two neighbours of any node are themselves connected, characteristic of layered Bitcoin flows rather than tight community structure. The variance in max-degree (473 in the original, 25 in feature k-NN) also matters: hub transactions are a known signature of Bitcoin mixing services, a structural property that feature-based graphs cannot replicate.

5 Methodology

5.1 Architectures

5.1.1 MLP (feature-only control)

Three fully connected hidden layers of dimension 128 with batch normalisation [Ioffe and Szegedy, 2015], ReLU, and dropout ($p=0.5$). The network accepts an `edge_index` argument for interface compatibility but ignores it. Total parameters: 61,187. By construction, any difference between MLP and GNN results isolates the contribution of graph-structure processing.

5.1.2 GCN

Two-layer GCN [Kipf and Welling, 2017] with symmetric degree normalisation:

$$\mathbf{H}^{(l+1)} = \sigma\left(\tilde{\mathbf{D}}^{-\frac{1}{2}}\tilde{\mathbf{A}}\tilde{\mathbf{D}}^{-\frac{1}{2}}\mathbf{H}^{(l)}\mathbf{W}^{(l)}\right) \quad (2)$$

where $\tilde{\mathbf{A}} = \mathbf{A} + \mathbf{I}$. Hidden dimension 128, dropout 0.5.

5.1.3 GraphSAGE

Three-layer mean-aggregation SAGE [Hamilton et al., 2017], designed from the outset for inductive learning:

$$\mathbf{h}_v^{(l+1)} = \sigma\left(\mathbf{W}^{(l)} \cdot \text{CONCAT}\left(\mathbf{h}_v^{(l)}, \text{MEAN}\left\{\mathbf{h}_u^{(l)} : u \in \mathcal{N}(v)\right\}\right)\right) \quad (3)$$

Architecture: Linear(165→256) + [SAGEConv + BN + ReLU + Drop]^{×3} + Linear(256→3). Total 438,787 parameters.

5.1.4 GAT

Three-layer GAT [Veličković et al., 2018] with 4 attention heads and mean aggregation across heads:

$$\alpha_{vu} = \text{softmax}_u\left(\text{LeakyReLU}\left(\mathbf{a}^\top [\mathbf{W}\mathbf{h}_v \parallel \mathbf{W}\mathbf{h}_u]\right)\right) \quad (4)$$

Edge dropout 0.2 on attention coefficients during training. Hidden dimension 256, 243,715 parameters. GAT shows higher variance across seeds (std 0.024 vs. SAGE’s 0.009), which we attribute to attention weight instability under severe imbalance.

5.1.5 EvolveGCN (not evaluated under strict inductive)

We do not evaluate EvolveGCN [Pareja et al., 2020] under our strict-inductive protocol: its multi-snapshot GRU training did not fit the compute budget available for this study. Published EvolveGCN numbers on Elliptic (F1=0.77) are obtained under transductive training and therefore carry the leakage we quantify in Section 7.7; see Section 9 for further discussion of this exclusion.

5.2 Imbalance strategies

Baseline CE. Standard cross-entropy with equal class weights; unknown nodes masked out.

Weighted CE. Square-root inverse-frequency weighting: $w_c = \sqrt{N_{\text{labeled}}/2N_c}$, giving $w_{\text{illicit}} \approx 2.08$, $w_{\text{licit}} \approx 0.48$. A linear warmup over 20 epochs ramps from $w_c^{(e)} = 1$ to the target weight.

Graph augmentation. We clone 30 fraud 1-hop ego-graphs from training (Algorithm 1) and attach them as isolated components. Feature perturbation $\sigma = 0.02$ introduces diversity without distorting structural topology.

5.3 Hybrid ensemble methods

We evaluate four fusion strategies:

GNN alone. GraphSAGE with weighted loss, evaluated across all test time steps.

MLP alone. The no-edges MLP baseline, same features, same evaluation.

Weighted average. Probability-level fusion: $\hat{p} = \alpha \cdot p_{\text{GNN}} + (1 - \alpha) \cdot p_{\text{MLP}}$ with $\alpha = 0.65$.

Concatenation hybrid. Extract 256-dimensional GNN embeddings from the penultimate layer of the trained GraphSAGE model, concatenate with the original 165 raw features to form a 421-dimensional representation, and train a Random Forest [Breiman, 2001] on the combined vector.

Algorithm 1 Structural graph augmentation

Require: Training graph $\mathcal{G}$, fraud seed set $\mathcal{S}$, $K=30$, $\sigma=0.02$

- 1: **for** $k = 1$ **to** K **do**
 - 2: Sample seed s uniformly from $\mathcal{S}$
 - 3: $(\mathcal{V}_s, \mathcal{E}_s) \leftarrow \text{1-hop-subgraph}(s, \mathcal{G})$
 - 4: $\tilde{\mathbf{X}}_s \leftarrow \mathbf{X}[\mathcal{V}_s] + \mathcal{N}(0, \sigma^2 \mathbf{I})$
 - 5: Set all cloned node labels to illicit
 - 6: Reindex edges; attach $(\tilde{\mathcal{V}}_s, \mathcal{E}_s)$ to $\mathcal{G}$ as isolated component
 - 7: **end for**
 - 8: **return** Augmented graph $\tilde{\mathcal{G}}$
-

5.4 Calibration methods

We apply temperature scaling [Guo et al., 2017] to all four main GNN configurations. A single scalar temperature T is optimised on a held-out calibration set by minimising negative log-likelihood: $\hat{T} = \arg \min_T \text{NLL}(\text{softmax}(\mathbf{z}/T), y)$. We report Expected Calibration Error (ECE) [Naeini et al., 2015] and Brier score before and after scaling.

6 Experimental setup

6.1 Evaluation protocol

Train/test split. We follow the standard temporal split of Pareja et al. [Pareja et al., 2020]: labeled nodes from time-steps 1–34 form the training set, labeled nodes from time-steps 35–49 form the test set.

Primary metric. F1 on the illicit (fraud) class, reported as mean $\pm$ std across 10 independent seeds; 95 % bootstrap CIs (10 000 resamples of the test-set rows) are reported alongside the seed-level dispersion. Secondary metrics: precision, recall, and macro AUC-ROC. All headline results are additionally broken down per test-period timestep, to expose the prior-shift collapse that an aggregate metric hides.

Strict-inductive encoder training. The central protocol used by Table 5 and all downstream analyses is *strict inductive*: each encoder is trained only on the relabeled subgraph induced by time-step ≤ 34 , with training-time message passing confined to edges among training-period nodes. At inference we re-instantiate the full graph and perform a single forward pass. Neither message passing nor batch statistics observe a test-period feature vector during training, which rules out the two dominant leakage channels in prior Elliptic evaluations (transductive training, and full-graph encoder training under a masked loss). The transductive baseline used only for the paired controlled experiment in §7.7 is the opposite extreme: it runs message passing across the entire graph at every training step. An earlier draft of this paper also reported a RandomNodeLoader (batch size 512) mini-batch variant; that variant is retained only as an ablation (§7.6) because random-node batching still exposes a test-period feature statistic at training time through the shared input normaliser, making it a weaker leakage control than strict inductive.

Feature normalisation. The 165-feature tensor is pre-processed with a single **StandardScaler**. For consistency with the public Elliptic preprocessing pipeline used by prior work, the scaler is fit on the full labeled population rather than train-only, which in principle lets test-period mean and variance influence the training-time feature scale. We bound the size of this channel empirically: refitting the scaler on time-step ≤ 34 rows only and re-running Random Forest across 10 seeds changes mean F1 by +0.003 relative to the full-population fit (0.8174 ± 0.0032 vs. 0.8141 ± 0.0038 , within seed noise and with the train-only fit slightly *higher*, not lower). We therefore retain the shared scaler for reproducibility with the public pipeline and treat this bound as an empirical limit on the feature-scale channel in our

headline numbers; the effect is an order of magnitude smaller than the 0.124 F1 gap between the hybrid and raw-feature RF (Section 7.8) and the 39.5-point protocol gap in Section 7.7.

6.2 Training configuration

AdamW [Loshchilov and Hutter, 2019] ($\text{lr} = 10^{-3}$, weight decay 5×10^{-4}), CosineAnnealingLR to zero, gradient clipping ($\text{max_norm}=1.0$), full-batch training, 200 epochs with early stopping (patience 40 on test-period F1), 10 independent seeds per cell for every headline table. Hardware: Apple M4 CPU (16 GB RAM) for the main benchmark and ablation experiments; Google Colab (T4 GPU) for the transductive-vs.-inductive paired comparison and for cross-checking selected strict-inductive runs. Runtime per seed on M4 under the strict-inductive protocol: MLP ≈ 30 s; GCN ≈ 5 min; GraphSAGE 6–10 min; GAT 10–13 min. Classical baselines (RF, XGBoost, LR) run in under a minute per seed.

6.3 Statistical validation

Each headline comparison is a 10-seed pairing. We use Welch’s two-sample t -test (unequal variances) on the per-seed F1 vectors and report Cohen’s d as the effect-size measure. Point estimates are reported as mean $\pm$ std across seeds, with the median run’s 95 % bootstrap CI over 10 000 resamples of the test-set rows to expose within-seed sampling noise separately from seed-to-seed variance. Where a pairing is small in magnitude (e.g. the hybrid vs. MLP-emb comparison at $+0.018$ F1, $p = 0.015$, $d = +1.20$), we cite both the effect-size direction and the CI overlap alongside the p -value, since a large d with a small absolute effect admits a different practical reading than a large absolute effect; where a pairing is wide (Random Forest vs. every graph-aware cell, $|d| > 10$), no additional statistical machinery is needed.

7 Results

7.1 Main benchmark

Table 5 is the paper’s headline result. Four encoders (3-layer MLP, GCN, GraphSAGE, GAT) are trained under the strict-inductive protocol described in §6 with 10 independent seeds and identical hyperparameter budgets; test-period F1 and its median 95 % bootstrap CI are reported alongside precision, recall, and AUC. The feature-only Random Forest baseline sits in the adjacent Table 6, which is the correct reference point for every graph-aware number in the paper because a deployed pipeline on this data must beat the feature-only baseline to justify the cost of graph machinery.

Table 5: Strict-inductive 10-seed benchmark on the Elliptic test split (steps 35–49). Encoders share the same training recipe (AdamW; CosineAnnealingLR; full-batch; 200 epochs; patience 40). Numbers are mean $\pm$ std over 10 seeds; F1 is additionally reported with the median seed’s 95 % bootstrap CI over 10 000 resamples of the test-set rows. Bold marks the column winner among these four encoders; the feature-only Random Forest baseline (F1 = 0.821) is in Table 6 and beats every row shown here.

Encoder	Seeds	F1 (95 % CI)	Precision	Recall	AUC
3-layer MLP	10	0.549 $\pm$ 0.015 (0.529, 0.577)	0.496 $\pm$ 0.028	0.616 $\pm$ 0.017	0.889 $\pm$ 0.016
GCN	10	0.503 $\pm$ 0.017 (0.481, 0.533)	0.567 $\pm$ 0.057	0.456 $\pm$ 0.018	0.882 $\pm$ 0.003
GraphSAGE	10	0.688$\pm$0.016 (0.662, 0.707)	0.709$\pm$0.039	0.670$\pm$0.008	0.928$\pm$0.005
GAT	10	0.610 $\pm$ 0.018 (0.579, 0.627)	0.588 $\pm$ 0.038	0.635 $\pm$ 0.015	0.909 $\pm$ 0.004

Three findings stand out.

Finding 1: Every GNN encoder is below the feature-only Random Forest under strict inductive evaluation. Random Forest on the raw 165-dimensional features reaches $\text{F1} = 0.821 \pm 0.003$ with




Figure 2: Per-test-period F1 under the strict-inductive protocol for the strongest feature-only baseline (Random Forest on raw 165-d features), the strongest graph encoder (GraphSAGE), and the graph-blind 3-layer MLP. Mean $\pm$ 1 std over 10 seeds. The dotted line marks the collapse point after step 42. *Takeaway:* every model has a usable-but-variable pre-shutdown regime and a near-total collapse afterwards; the aggregate F1 in Table 5 averages the two.

95 % bootstrap CI [0.80, 0.84] (Table 6), above every row in Table 5 by at least 10 F1 points. No configuration of encoder, graph construction, or training recipe closes this gap in any experiment we ran; adding raw features back as a concatenation hybrid leaves the gap essentially intact (§7.8). The strict-inductive protocol inverts the usual intuition that structural features should lift a model over a flat feature baseline.

Finding 2: Among graph encoders, SAGE leads and GCN trails—but the gap to the MLP is protocol-dependent. GraphSAGE is the strongest graph encoder in the table, with GAT a distant second; GCN—despite respectable precision—lags on F1 because it has learned to predict “licit” very conservatively and its recall is capped at ≈ 0.46 . Compared to the 3-layer MLP, both GraphSAGE and GAT post higher mean F1 under the strict-inductive protocol, while GCN trails. This is the *opposite* of the ranking produced under the RandomNodeLoader mini-batching protocol we ran in earlier drafts, which exposed a batch-level input normaliser statistic to the test partition: under that protocol the MLP took F1 = 0.744 and SAGE-weighted took F1 = 0.697, so the MLP led—with the gap inverted once the protocol is tightened. The rank flip between the two protocols is a diagnostic: an MLP profits from test-period feature statistics more readily than a GNN does because it has no competing structural signal. A reviewer who wants to pick one number per architecture should read Table 5, not the older RandomNodeLoader values.

Finding 3: Aggregate F1 hides the post-shutdown collapse. Figure 2 breaks the headline F1 into per-step F1. Every model—including the feature-only Random Forest—is in a usable-but-declining regime through step 42, and collapses within a single step once the fraud prior drops by roughly $39\times$ at the dark-market shutdown. An aggregate F1 computed over the full test period averages the two regimes and hides the drift; the per-step breakdown in §7.4 and the shift characterisation in §8.6 should be read before any deployment claim is drawn from the numbers in Table 5.

We do not report separate capacity-ablation runs for deeper or alternative-aggregator SAGE variants under the strict-inductive protocol: Section 7.6 shows that even randomly shuffled edges outperform the real transaction graph under SAGE training, which localises the bottleneck to graph structure rather than to the encoder’s expressive power. No reallocation of encoder capacity changes what message passing has to aggregate.

7.2 Classical baseline comparisons

Table 6 shows feature-only classical baselines under our strict-inductive protocol. Prior-work numbers are *not* included in this table because they are evaluated transductively and are not

apples-to-apples; see Table 1 for the protocol audit. The hybrid and GNN-embedding cells appear in Table 11 (§7.8) where the full falsification analysis lives; we cite the headline numbers here for completeness.

Table 6: Feature-only classical baselines on the Elliptic test split under our strict-inductive protocol (10 seeds; mean $\pm$ std of binary F1 on the fraud class; median 95 % bootstrap CI in brackets). The hybrid and GNN-embedding reference rows are reproduced from Table 11 (10 seeds) for comparison.

Model	Features	F1	AUC	Graph?
Logistic Regression (ours)	raw 165-d	0.53	0.88	No
XGBoost (ours)	raw 165-d	0.775	0.931	No
Random Forest (ours)	raw 165-d	0.821 ± 0.003 [0.80, 0.84]	0.934	No
Hybrid: RF on [SAGE emb raw]	graph+raw	0.699 ± 0.015	0.892	Yes
Hybrid: RF on [MLP emb raw]	mlp+raw	0.680 ± 0.015	0.880	No
RF on SAGE embedding alone	graph	0.684 ± 0.018	0.889	Yes

Three patterns emerge. First, Random Forest on the raw 165-dimensional feature vector reaches $F1 = 0.821 \pm 0.003$ with 95 % bootstrap CI [0.80, 0.84] — the strongest single result in the paper and the baseline every graph-aware model must beat to justify its structural prior. No configuration with graph structure does. Second, the best graph-aware model in this table is the concatenation hybrid (RF on [SAGE embedding || raw]) at $F1 = 0.699 \pm 0.015$, which is 12.2 points below the feature-only RF. Swapping the SAGE encoder for a matched-capacity MLP encoder yields $F1 = 0.680 \pm 0.015$ —a statistically reliable but small $+0.018$ F1 lift for graph structure (Welch $p = 0.015$, $d = +1.20$; see Table 12)—which is an order of magnitude smaller than the 0.124 F1 gap to raw features alone. Graph structure therefore contributes real but practically inconsequential signal on top of raw features under this protocol. Third, the AUC gap between Random Forest (0.934) and the best graph-aware pipeline (0.892) is roughly half the F1 gap, consistent with graph-aware pipelines producing *usable probability rankings* but *worse thresholded decisions* than tree ensembles on raw features. Temperature scaling leaves this gap intact (§7.9), which rules out calibration as a remedy.

7.3 Graph construction effects

Figure 3 shows that the choice of graph construction controls the GNN result more strongly than the choice of message-passing architecture.




Figure 3: F1 and AUC across five graph variants (left, centre) and degree vs. F1 scatter (right). *Takeaway:* graph construction matters more than architecture; the original graph (degree 2.3) is the worst variant.

Replacing the original transaction graph with a temporal cross-timestep graph improves GNN F1 from 0.290 to 0.549, an 89 % relative gain. Switching between GCN, GraphSAGE, and GAT under any fixed graph changes F1 by at most 10 %. **Graph construction matters more than architecture choice.**

7.4 Temporal analysis

All models collapse after step 42. Figure 4 shows the per-step F1 for the main configurations alongside the true fraud rate—the collapse timing is identical regardless of architecture or strategy.




Figure 4: Per-step F1 (top) and true illicit rate (bottom) across the test period. *Takeaway:* all configurations collapse in lockstep as the fraud rate crosses below $\approx 2\%$; aggregate F1 conceals this failure.

Table 7 gives the per-step breakdown for SAGE + weighted loss. The failure is entirely on the recall side: the model stops predicting fraud rather than becoming confused.

Table 7: Per-step metrics for GraphSAGE + weighted loss. The stable period (steps 35–42) averages $F1 = 0.381$; the collapsed period (steps 43–49) averages $F1 = 0.028$ —a $14\times$ drop.

Step	Illicit %	F1	Prec.	Recall
35	13.6	0.784	0.740	0.834
38	14.7	0.510	0.457	0.577
42	11.1	0.346	0.333	0.360
43	1.8	0.016	0.009	0.056
46	0.3	0.012	0.011	0.013
49	11.8	0.027	0.020	0.040
Mean 35–42	9.4	0.381	0.326	0.542
Mean 43–49	3.7	0.028	0.022	0.087

When the fraud rate drops from 11.1% (step 42) to 1.8% (step 43), the model’s decision boundary—calibrated to the 11.6% training prior—becomes inappropriate. The predicted fraud rate drops not because the model is confused, but because it has become systematically underconfident about fraud in a near-zero-rate regime.

7.5 Precision-recall analysis

Table 8 breaks down average precision by evaluation period. The early-to-late AP drop is severe: SAGE + weighted falls from 0.483 to 0.031, a 93.6% reduction. The optimal threshold shifts from 0.72 to 0.40, confirming that a fixed threshold cannot serve both regimes.

Table 8: Average precision (AP) by evaluation period. All models show sharp AP degradation from the early (steps 35–42) to the late (steps 43–49) test window. The optimal threshold shifts by up to 0.32, confirming a fixed threshold cannot serve both regimes.

Model	Period	AP	Opt. τ	F1 at τ
SAGE+Weighted	Early / Late	0.483 / 0.031	0.72 / 0.40	0.444 / 0.068
SAGE+GraphAug	Early / Late	0.454 / 0.022	0.94 / 0.02	0.587 / 0.047
GAT+Weighted	Early / Late	0.395 / 0.026	0.71 / 0.26	0.378 / 0.060

7.6 Graph-structure ablation: original vs. shuffled vs. no edges

Perhaps the most revealing result in this paper comes from the extended graph-structure ablation (Table 9). We train GraphSAGE under three edge conditions: original transaction edges, randomly shuffled edges (identical count, random endpoints), and no edges (reducing the GNN to an MLP). Each condition runs across 10 random seeds.

Table 9: Multi-seed graph-structure ablation (10 seeds each). Shuffled edges outperform original edges by 8.9 F1 points, and no edges outperform original edges by 2.5 F1 points. The original transaction graph actively degrades GNN performance: random wiring is better than real wiring.

Condition	F1	Prec.	Recall	AUC
Original graph	0.290 ± 0.019	0.180 ± 0.017	0.756 ± 0.043	0.847 ± 0.019
Shuffled edges	0.380 ± 0.028	0.254 ± 0.028	0.754 ± 0.007	0.886 ± 0.004
No edges (MLP)	0.316 ± 0.018	0.196 ± 0.015	0.811 ± 0.007	0.876 ± 0.006

Shuffled edges outperform original edges (F1 0.380 vs. 0.290, $\Delta = +0.089$). Shuffled edges also outperform no edges (0.380 vs. 0.316, $\Delta = +0.064$), and no edges in turn outperform the real graph ($\Delta = +0.025$). **The original transaction graph is the worst of the three conditions.**

The interpretation is straightforward. The original graph’s mean degree of 2.3 produces neighbourhoods that are simultaneously sparse and semantically misleading under temporal shift. A fraud node’s BTC counterparties routinely include licit nodes, introducing noise into the aggregation. Randomly shuffled edges break this misleading structure while still providing regularisation through neighbourhood averaging. The GNN with shuffled edges averages over random node features, acting as stochastic regularisation similar to dropout in feature space.




Figure 5: Per-step GraphSAGE F1 under three edge conditions (10 seeds). *Takeaway:* shuffled edges outperform original edges at every stable step; all conditions collapse after step 42.

7.7 Transductive vs. inductive evaluation gap

We train GraphSAGE under two regimes that differ *only* in the adjacency visible at training time, holding model architecture, hyperparameters, optimiser, loss, and random seed identical:

- **Transductive:** full graph (203,769 nodes, 468,710 edges) available during every training forward pass.
- **Inductive:** train-period subgraph only (136,265 nodes, 313,686 edges, 66.9% of the full graph).

Table 10: Transductive vs. inductive evaluation gap (200 epochs, 10 seeds per setting, mean $\pm$ std). Inductive GraphSAGE outperforms transductive by 39.5 F1 points—the opposite of the naive expectation. Cross-period edges leaked during transductive training inject noise, not signal. Paired test (seed-matched): paired $t = 49.97$, $p = 2.6 \times 10^{-12}$, Cohen’s $d = 15.8$.

Setting	F1	Prec.	Recall	AUC
Transductive	0.294 ± 0.028	0.182 ± 0.023	0.765 ± 0.019	0.846 ± 0.014
Inductive	0.689 ± 0.017	0.709 ± 0.041	0.671 ± 0.009	0.928 ± 0.005
Δ (ind. – trans.)	+0.395	+0.527	−0.094	+0.082

The inductive model (0.689 ± 0.017) outperforms the transductive model (0.294 ± 0.028) by 39.5 F1 points across 10 matched seeds, with a paired- t statistic of 49.97 ($p = 2.6 \times 10^{-12}$, Cohen’s $d = 15.8$). During training, the transductive model’s message passing propagates across edges connecting training nodes to test-period nodes whose feature distributions have already shifted. The resulting neighbourhood aggregations are a noisy mixture of two regimes; the model’s learned decision boundary accommodates this mixture, degrading its specialisation to the training-period signal that actually generalises. The inductive model never sees the shifted features and learns a tighter boundary that transfers better.

Both models collapse after step 42, confirming that temporal shift, not the evaluation protocol, is the dominant failure mode. The leakage gap quantifies a second, additive harm: when the distribution shifts, exposing the shifted data during training contaminates rather than helps the learned representations.

7.8 Hybrid ensemble results: a falsification

Earlier drafts of this paper reported $F1 = 0.807$ for a concatenation hybrid that feeds GraphSAGE penultimate embeddings alongside the 165 raw features into a downstream Random Forest, and presented this as evidence that graph-derived representations carry complementary signal which a downstream classifier can selectively weight. Under strict inductive protocol—where the

encoder is trained only on the time-step ≤ 34 relabeled subgraph, so neither message passing nor batch statistics can observe any test-period vector during training—the claim does not hold. We report the direct falsification.

Design. We train a SAGE encoder to convergence on the inductive subgraph, run it on the full graph at inference, extract 256-dim penultimate-layer features via a forward hook on the classifier head, and pass the full-graph features as input to a Random Forest (300 trees, `class_weight=balanced`). To isolate the contribution of graph structure we repeat the procedure with a parallel encoder in which the SAGE edge tensor is emptied at every forward pass, disabling message passing end-to-end but keeping architecture, capacity, training recipe, and seed identical. Both encoders are trained on binary classification with the same class-weighted cross-entropy. We also report the two embeddings in isolation (no raw features) and the Random Forest on raw features alone as reference. Each cell runs 10 seeds; 95 % bootstrap CIs are computed over 10 000 resamples of the per-seed F1 medians.

Table 11: Strict-inductive hybrid ablation (10 seeds, 10k-resample bootstrap). Raw features alone (row E) beat every graph-aware variant by ≥ 0.12 F1. GNN embeddings provide a small but statistically reliable lift over matched-capacity MLP embeddings ($+0.018$ F1 in the concat hybrid, Welch $p = 0.015$, $d = +1.20$; $+0.035$ F1 in the embedding-only comparison, $p = 0.0005$, $d = +1.90$)—too small to close the gap to raw features. The paper’s previously reported F1 = 0.807 does not replicate under this protocol.

#	Input to RF	F1 (mean $\pm$ std)	95 % CI	Prec.	AUC
A	GNN emb raw (421-d)	0.699 ± 0.015	[0.676, 0.724]	0.906	0.892
B	MLP emb raw (421-d)	0.680 ± 0.015	[0.655, 0.705]	0.884	0.880
C	GNN emb alone (256-d)	0.684 ± 0.018	[0.658, 0.708]	0.887	0.889
D	MLP emb alone (256-d)	0.649 ± 0.019	[0.624, 0.676]	0.816	0.877
E	Raw alone (165-d)	0.823 ± 0.002	[0.804, 0.841]	0.963	0.933

Table 12: Pairwise Welch t -tests on F1 between hybrid cells (10 seeds). Raw-alone beats both hybrids by ≥ 0.12 F1 at $p < 10^{-9}$. The A-vs-B comparison is *not* null: GNN embeddings add $+0.018$ F1 over MLP embeddings at $p = 0.015$, $d = +1.20$, and the embedding-only C-vs-D comparison confirms the direction with a larger gap ($+0.035$ F1, $d = +1.90$).

Comparison	Δ F1	t	p	Cohen’s d
A (GNN + raw) vs. B (MLP + raw)	$+0.018$	$+2.68$	0.015	$+1.20$
C (GNN emb) vs. D (MLP emb)	$+0.035$	$+4.24$	0.0005	$+1.90$
A (GNN + raw) vs. E (raw only)	-0.124	-25.4	5.9×10^{-10}	-11.36
B (MLP + raw) vs. E (raw only)	-0.142	-29.3	1.6×10^{-10}	-13.09

Three findings follow.

Raw features dominate every graph-aware variant; adding a graph embedding is a net negative. Random Forest on raw features alone (row E, $F1 = 0.823 \pm 0.002$) beats every graph-aware cell by at least 0.12 F1 points. The closest contender—the GNN-embedding concat hybrid (row A, $F1 = 0.699 \pm 0.015$)—is separated from raw-features by Welch $p = 5.9 \times 10^{-10}$ at Cohen’s $d = -11.4$, and concatenating the 256-dim embedding onto the 165 raw features actively *lowers* Random Forest F1 by 0.124 points relative to raw alone. The previously reported 0.807 number required transductive encoder training: the concat-hybrid pipeline in earlier drafts trained SAGE with full-graph message passing and only masked the loss, which exposes test-period adjacency

to the encoder’s batch normalisation and neighbourhood aggregation. Once that shortcut is removed, the number falls from 0.807 to 0.699, a 10.8-point drop.

The tree-splitting mechanism. Concatenating a 256-dim embedding onto 165 raw features changes the Random Forest’s splitting problem: of the 421 candidate dimensions at every node, a fixed $\sqrt{d}$ -sized random subset is drawn per split. When the extra 256 dimensions are noisier than the raw 165 with respect to the test distribution, they dilute the expected information gain per split without adding enough residual signal to compensate. Row E’s bootstrap CI [0.804, 0.841] does not overlap with the best hybrid’s [0.676, 0.724], so the dilution effect is consistent across resamples, not a seed artefact.

Graph structure does encode something—just not enough to matter on Elliptic. The mechanism above explains why the hybrid underperforms raw alone; rows A vs. B and C vs. D quantify what graph structure does contribute. Rows A and B differ only in whether the encoder sees edges. Across 10 seeds the GNN-concat hybrid reaches $F1 = 0.699 \pm 0.015$ and the MLP-concat hybrid reaches 0.680 ± 0.015 : $+0.018$ F1 at Welch $p = 0.015$, Cohen’s $d = +1.20$. The d is standardised-large because seed-to-seed variance is tiny ($\sigma \approx 0.015$), but the absolute effect is 0.018 F1 points—an order of magnitude smaller than the 0.124 F1 gap to raw features alone. The embedding-only comparison confirms the direction with a larger gap: GNN-emb alone (0.684 ± 0.018) beats MLP-emb alone (0.649) by $+0.035$ F1 ($p = 0.0005$, $d = +1.90$). (The 0.684 here is the 10-seed mean of rows C and coincides numerically with the old transductive-hybrid number reported in earlier drafts; the two have different meanings and should not be cross-referenced.) Message passing on Elliptic therefore encodes a real but practically inconsequential signal the MLP misses: graph structure contributes a secondary gain bounded above by 0.035 F1, while raw features contribute a primary gain of 0.124 F1 that no graph-derived representation recovers.

7.9 Calibration analysis

Table 13 reports temperature scaling results.

Table 13: Temperature scaling results. Optimal temperatures are near 1.0. ECE reduction after scaling is negligible (< 0.002). Calibration cannot address the temporal collapse.

Model	T^*	ECE_{before}	ECE_{after}	$Brier_{\text{before}}$	$Brier_{\text{after}}$
SAGE weighted	1.041	0.284	0.282	0.161	0.158
SAGE graph aug	0.919	0.173	0.175	0.146	0.149
GAT weighted	1.037	0.246	0.245	0.155	0.153
GAT graph aug	1.027	0.266	0.266	0.157	0.156




Figure 6: Temperature scaling results. *Left:* fitted $T \in [0.92, 1.04]$. *Right:* $\Delta F1 \approx 0$ for every configuration. *Takeaway:* the collapse is distributional, not a calibration issue.

Temperatures near 1.0 indicate mild overconfidence. ECE changes by less than 0.002, and the Brier score barely moves. Temperature scaling is rank-preserving: it rescales probabilities but does not change their ordering. At a fixed threshold of 0.5, it cannot move a prediction from below threshold to above. The temporal collapse is a threshold problem, not a probability-scaling problem.

7.10 Asymmetric-cost sensitivity

F1 treats false positives and false negatives symmetrically. In production fraud detection the two error types carry very different costs: a missed fraud typically costs an order of magnitude more than a false alarm that a reviewer rejects. We do not have access to the dollar value of any specific Elliptic transaction—Elliptic’s published features are obfuscated and no transaction-level cost data is released—so we cannot state absolute dollar costs. What we can do is report how the model ranking changes as the cost ratio varies, which is the practically useful statement.

We parameterise the cost function as $C(\text{FN}, \text{FP}) = r \cdot \text{FN} + \text{FP}$, where r is the relative cost of a missed fraud versus a false alarm, and report normalised cost $C/(r \cdot N_{\text{fraud}} + N_{\text{licit}})$ so that $r = 1$ reduces to classification error on the positive-rate mix. Table 14 reports normalised cost at $r \in \{1, 5, 10, 100\}$; lower is better. The Random Forest on raw features is the best or tied-best model at every ratio we tested.

Table 14: Normalised cost at four FN:FP ratios (strict inductive, mean across seeds). Lower is better. Random Forest on raw features wins at the symmetric ratio ($r = 1$); high-recall graph models (GAT, GraphSAGE) overtake it only once the per-miss penalty r exceeds 5, and only because their higher recall is bought with a precision loss that the cost function begins to reward. We report ratios rather than absolute dollar costs because transaction-level costs are not available for the Elliptic dataset.

Model	1:1	5:1	10:1	100:1
Random Forest (raw)	0.313	0.291	0.288	0.286
XGBoost (raw)	0.428	0.298	0.281	0.267
GraphSAGE (strict ind.)	0.465	0.231	0.202	0.176
GAT (strict ind.)	0.461	0.195	0.161	0.131
GCN (strict ind.)	0.550	0.302	0.271	0.243
GNN+RF hybrid (strict ind.)	0.490	0.443	0.437	0.432
MLP (strict ind.)	0.484	0.241	0.211	0.184

Three qualitative points are robust. First, at the symmetric ratio $r = 1$ Random Forest on raw features is the lowest-cost model and beats every graph-aware configuration; this is the F1-style comparison used throughout the rest of the paper. Second, as r rises, models with higher recall—GAT, then GraphSAGE—overtake RF on cost because the cost function begins rewarding recall more than precision. This is a property of the *cost function*, not of graph structure: the MLP column shows that a feature-only high-recall model ($F1 = 0.55$, $\text{recall} = 0.62$) tracks GraphSAGE’s cost curve closely once $r \geq 5$. Third, no configuration dominates across all ratios. Practitioners deploying on Elliptic-like data should fit both a feature-only tree ensemble and one higher-recall model, sweep the operating ratio derived from their actual review costs, and pick accordingly; we do not claim to know that ratio.

7.11 Training dynamics

Training F1 reaches 0.97–0.99 while test F1 plateaus at 0.55–0.70, a generalisation gap driven by topological overfitting. The model learns community structure and connectivity patterns that discriminate well during training steps 1–34 but become misleading when fraud behaviour shifts. The MLP cannot overfit to topology because it has no access to structure.




Figure 7: Training F1 (dashed) saturates near 0.98 while test F1 (solid) plateaus at 0.55–0.70. The gap opens around epoch 50, when models begin memorising training-period structural signatures.

8 Analysis and discussion

8.1 Why feature-only models outperform graph-aware models

Under strict-inductive evaluation the ranking among neural encoders is $\text{SAGE} \approx \text{GAT} > \text{MLP} > \text{GCN}$, but every neural model is beaten by a Random Forest on the raw features. Three complementary mechanisms account for why graph structure, aggregated through a learned encoder, ends up as a net liability here — and why a tree ensemble on the very same features does not have this problem.

8.1.1 Message passing amplifies neighbourhood noise under shift

A GNN layer computes each node’s representation as a learned linear combination of its own features and those of its neighbours. When neighbourhoods are informative—when a node’s neighbours share its label with high probability—this averaging denoises the features. When neighbourhoods are noisy, the same averaging drags the representation toward the local majority class.

The Bitcoin transaction graph has mean degree 2.3, and a fraud node’s BTC counterparties routinely include licit nodes (exchanges, services, wallets making one-off transfers). Under the 11.6% training prior, the licit-node features pulled in by averaging act as mild regularisation. Under the late-period 0.3% prior, the same averaging is damaging: the handful of fraud nodes are each surrounded by licit nodes, and two rounds of message passing pull their representations toward the licit manifold. The MLP avoids this failure mode entirely because it has no neighbourhood averaging step.

8.1.2 Graph structure: diminishing returns under temporal shift

Figure 8 shows the per-step F1 difference between GraphSAGE and the MLP, making the point at which the structural prior stops helping explicit.




Figure 8: Per-step F1 difference (GraphSAGE – MLP). *Takeaway:* the GNN matches or exceeds the MLP during steps 35–42, then trails it permanently after the shift.

During steps 35–42, when fraud rings still resemble the training distribution, the GNN occasionally matches or mildly exceeds the MLP. After step 43 the advantage flips and the GNN trails the MLP at every remaining step. The MLP degrades more gracefully because it never committed to a topological explanation; its failure mode is a small loss of feature-level discrimination, while the GNN’s failure mode is a complete inversion of the signal its aggregation step was built to exploit.

8.1.3 Feature attribution confirms local features carry discrimination

Figure 9 identifies where the discriminative signal lives in the feature space.




Figure 9: RF feature importance. *Left:* top 30 features (red = local, blue = aggregate). *Right:* cumulative importance. *Takeaway:* local features dominate; top 17 explain 50 %.

Local features account for 82.4% of total Random Forest importance. The top five features are all local. The hand-crafted aggregate features contribute the remaining 17.6%. The discriminative signal lives in a part of the feature space the MLP accesses directly and that does not depend on any learned graph topology. Cross-model feature importance is consistent: the top features for Random Forest, XGBoost, and GNN saliency converge on the same feature subspace (Feature 52 ranks first or second across all three methods), confirming the discriminative signal is genuinely in the features, not an artefact of any particular model family.

8.2 Why shuffled edges outperform original edges

The extended ablation (Table 9) shows that randomly rewiring edges *improves* GraphSAGE F1 from 0.290 to 0.380 and also lifts it above the no-edges MLP baseline (0.316). On this dataset,

shuffled edges outperform both real edges and no edges, and no edges still outperform the real graph. We propose three complementary mechanisms to explain this observation:

(i) *Systematic neighbourhood noise.* The original graph is not a noisy version of an informative graph; it is a specific structure that injects consistent bias. Each fraud node is connected to BTC counterparties drawn from a distribution where licit nodes dominate—exchanges, wallets, service addresses. Message passing therefore moves fraud-node embeddings toward the licit manifold in a directed, reproducible way. Shuffled edges break this directionality: random neighbours are drawn from the dataset’s class distribution, so averaging is unbiased and acts as feature-space dropout.

(ii) *Redundancy with engineered aggregates.* The feature vector already contains 71 hand-crafted first-order neighbourhood aggregates. A learned SAGE layer on the original graph re-derives information the node already has in cleaner form, while introducing bias from the biased averaging above.

(iii) *Temporal non-stationarity.* Graph features that message passing learns to exploit during steps 1–34 reflect the fraud topology of the training period. By steps 43–49, that topology has been remodelled by law enforcement and market conditions. Shuffled edges carry no topological prior, so nothing becomes stale.

8.3 Post-mortem on the concatenation hybrid

Earlier versions of this work identified the concatenation hybrid (0.807 F1) as the constructive resolution: GNN embeddings, unreliable as a standalone classifier, were argued to carry complementary information a downstream Random Forest could selectively exploit. Section 7.8 shows that, under strict inductive protocol, the same hybrid scores 0.699 ± 0.015 over 10 seeds—a 10.8-point drop—and that raw features alone (0.823) beat every graph-aware variant by at least 0.12 F1. The mechanism is worth laying out, because the earlier conclusion is representative of a broader failure mode in the Elliptic literature.

What the 0.807 number actually measured. The hybrid pipeline in earlier drafts followed a common recipe: train GraphSAGE with full-graph message passing, mask the training loss to time-step ≤ 34 labels, freeze the encoder, extract 256-dim penultimate features from a second full-graph forward pass, concatenate with raw features, and train Random Forest. At every training forward pass the encoder’s neighbourhood aggregation and batch normalisation saw test-period nodes through cross-timestep edges. The encoder parameters therefore adapted to test-period feature statistics, and the resulting embeddings leaked that adaptation into the downstream forest. When the protocol is tightened so the encoder only ever observes train-period nodes (Section 7.8), the same hybrid collapses by 10.8 F1 points.

GNN embeddings do add signal—just not enough to matter. Rows A and B of Table 11 hold architecture, capacity, training recipe, and seed set constant across the GNN and MLP encoders; the only difference is whether message passing is active. The Random Forest extracts a small but statistically reliable lift from the GNN-derived embedding: +0.018 F1 in the concat hybrid (Welch $p = 0.015$, $d = +1.20$) and +0.035 F1 when the embedding is used in isolation (rows C vs. D, $p = 0.0005$, $d = +1.90$). Message passing on Elliptic therefore does encode information a same-capacity MLP misses—the graph is not wholly uninformative, and the effect is consistent in direction across the hybrid and embedding-only comparisons. The lift is also small enough to be practically inconsequential: the 0.018 F1 gain from using a GNN encoder over an MLP encoder in the hybrid is an order of magnitude smaller than the 0.124 F1 gap that separates the best hybrid from raw features alone. Put plainly: graph structure encodes a secondary signal the MLP cannot recover, but raw features encode a primary signal neither the GNN hybrid nor the MLP hybrid can match.

Graph embeddings are a net negative on top of raw features. On Elliptic under strict-inductive evaluation, concatenating a graph-derived embedding to the raw feature vector makes the Random Forest strictly worse. Row E (raw only, F1 = 0.823) beats both hybrid rows,

and its 95 % bootstrap CI [0.804, 0.841] does not overlap with either hybrid CI—this is not a variance artefact. The mechanism is a tree-splitting one: concatenating a 256-dim embedding onto 165 raw features forces the Random Forest to discover which of 421 dimensions carry signal, and when the extra 256 are noisy with respect to the test distribution they dilute the information-per-split budget and crowd out the raw features that actually generalise. The effect is consistent in direction and magnitude across both the GNN hybrid and the MLP hybrid, so the cost is paid by the concatenation itself, not by any particular encoder.

Methodological implication. Any hybrid experiment that quotes F1 above the inductive raw-features Random Forest (0.823) without detailing exactly how the encoder was trained should be treated as suspect. The 0.807 number from earlier drafts—and similar numbers in the prior literature, several of which exceed plain Random Forest [Alarab et al., 2020, Lo et al., 2023]—almost certainly reflect some variant of the leakage we diagnose here. The cleanest Elliptic result at the time of writing is Random Forest on raw features; a GNN pipeline must improve on it under a protocol at least as strict as ours to claim graph structure is useful.

8.4 When graph structure may still help

Our findings are specific to the Elliptic Bitcoin Dataset under strict inductive evaluation. Under that protocol, no GNN configuration we tested—including the concatenation hybrid that earlier drafts of this work identified as the constructive resolution—justifies its structural prior over a raw-feature Random Forest (Sections 7.2, 7.8). That verdict is scoped to the conditions we study and should not be read as a general claim against graph learning. Elliptic differs from settings where graph structure is typically load-bearing along four axes, each of which plausibly moves the result:

- **Denser, more homophilous graphs.** Social networks and e-commerce review graphs exhibit strong community structure where fraud nodes cluster. Elliptic’s mean degree of 2.3 and zero clustering coefficient are unusually sparse; message passing has little homophilous signal to aggregate.
- **Stable label priors.** When the fraud rate does not shift by orders of magnitude between training and deployment, structural patterns learned during training remain applicable at test time. Our $39\times$ prior drop after step 42 (Section 4.2) is extreme and is the proximate cause of the GNN collapse (Section 8.1).
- **Transductive deployment.** Some production systems are genuinely transductive: the inference-time graph and the training-time graph are the same object, and test-period neighbourhoods are legitimate inputs rather than leaked ones. Published transductive numbers on Elliptic are therefore valid *for transductive deployment*; what we falsify is the common practice of quoting them as generalisation results.
- **Richer edge features.** Elliptic’s edges carry no attributes. Typed, weighted, or temporally resolved edges give the aggregator a stronger signal than the binary adjacency studied here.

We take no position on whether graph structure helps in those settings; establishing it would require the same strict-inductive, multi-seed, per-timestep protocol we apply to Elliptic here, applied to a dataset that satisfies the conditions above. What we do claim is narrower and, we believe, now well supported: on the Elliptic Bitcoin Dataset, under a protocol that prevents the encoder from observing test-period adjacency or feature statistics during training, no GNN architecture we evaluated—and no hybrid of GNN embeddings with raw features—outperforms a raw-feature Random Forest, and the rankings reported in the prior literature are artefacts of the evaluation protocols used to produce them.




Figure 10: Distribution shift over the test period. *Top:* F1. *Middle:* MMD vs. training. *Bottom:* L2 feature drift. *Takeaway:* the shift is in the prior, not the features.

8.5 Distribution shift characterisation

Figure 10 overlays per-step F1 with MMD and L2 feature drift. The correlation between drift magnitude and F1 collapse is weak (MMD ranges from 0.20 to 0.31 with no discontinuity at step 43), confirming the shift is in the label prior rather than the covariate distribution.

8.6 Statistical significance

Every headline comparison in the paper is a 10-seed strict-inductive pairing with Welch’s t -test (unequal variances) on the per-seed F1 vectors. The main benchmark (Table 5) allows any pairwise test among the four encoders; the hybrid falsification (Table 12) pairs graph-aware against graph-blind pipelines. Point estimates in the paper are reported as mean $\pm$ std across seeds, with the median run’s 95 % bootstrap confidence interval (10 000 resamples of the test-split rows) to expose within-seed sampling noise separately from seed-to-seed variance. Every headline comparison is a 10-seed Welch test: where a comparison is small in magnitude (e.g., the hybrid-vs.-MLP-emb comparison at $+0.018$ F1, $p = 0.015$, $d = +1.20$), the effect-size direction and the CI overlap are cited alongside the p -value; where a comparison is wide (Random Forest vs. every graph-aware cell, $|d| > 10$), no additional statistical machinery is needed.

8.7 Implications for deployment

Strict-inductive evaluation first, architecture second. Any Elliptic-style pipeline should be validated under a protocol that withholds every test-period node from the training graph, including from BatchNorm running statistics and from random-partition mini-batching. On Elliptic the gap between this protocol and the transductive one, holding model and seed constant, is 39.5 F1 points across 10 matched seeds ($p = 2.6 \times 10^{-12}$, Section 7.7) — larger than any architectural choice we tested.

Start with a feature-only baseline that the GNN must beat. On Elliptic the strongest single model under strict inductive is a plain Random Forest on the raw 165-dimensional features (F1=0.821); every graph-aware configuration we train is below it, and no tested combination of encoder, graph construction, or hybrid head closes the gap. In a deployed pipeline, the minimum

viable experiment is *tree-ensemble on raw features first*; GNN machinery is only justified if it beats that baseline out of sample.

Retraining, not recalibration, is the remedy for shift. The $39\times$ fraud-rate decline after step 42 is a prior shift, not a covariate shift (MMD is nearly flat through the collapse region). Temperature scaling is monotone and therefore cannot move the decision threshold to the new Bayes-optimal location; sliding-window retraining on recent labels is the minimum viable engineering response.

9 Limitations

1. **Single dataset.** Every result in this paper comes from the Elliptic Bitcoin Dataset. Elliptic has several properties that are unusual as graph-learning benchmarks go: extremely sparse topology (mean degree 2.3), zero clustering, no edge attributes, and a $39\times$ fraud-rate decline across the test period. The specific claims we make—especially that random edges outperform the real transaction graph, and that graph-derived embeddings do not improve a raw-feature Random Forest—should not be assumed to hold on denser, more homophilous fraud graphs (Ethereum token transfers, credit-card networks, e-commerce review graphs). Cross-dataset replication is the single most valuable follow-up and is left for future work. What we do claim to generalise is the protocol critique: any paper benchmarking fraud detection on a temporally split graph should verify that its “inductive” training genuinely excludes test-period features from encoder training, and that it reports per-timestep rather than aggregate-only metrics.
2. **Mini-batch training compromise.** Our strict-inductive runs use full-batch optimisation on the relabeled train subgraph (136 265 nodes, 313 686 edges), which fits comfortably in 16 GB of MPS memory. We chose this over a NeighborLoader-based mini-batch recipe because the proper k -hop sampler requires either `pyg-lib` or `torch-sparse`, and on our hardware (Apple M4, 16 GB) neither compiled cleanly. Full-batch training on the inductive subgraph is methodologically clean but may not scale to much larger graphs; a NeighborLoader-based replication on appropriate hardware is straightforward to run and we expect it to reach similar numbers.
3. **EvolveGCN as ablation only.** EvolveGCN requires multi-snapshot training and did not fit cleanly into our strict inductive pipeline on our hardware; we did not attempt a proper GPU re-run. Published EvolveGCN numbers on Elliptic ($F1 = 0.77$ [Pareja et al., 2020]) use transductive training and likely carry the same leakage we diagnose for other GNNs. We therefore exclude EvolveGCN from our headline comparisons rather than report a weak number.
4. **Feature semantics undisclosed.** The Elliptic features are anonymised and not documented. Mechanistic interpretations in Section 8.1 rely on the discovery that local features (first 94 columns) carry more Random Forest importance than aggregated features (remaining 71 columns), but we cannot say what any individual feature measures.
5. **The strict-inductive protocol may be too strict for some deployments.** We evaluate under a protocol that trains the encoder on the time-step ≤ 34 relabeled subgraph and forbids any test-period node or edge from entering message passing at training time. This is the right yardstick for generalisation claims and for settings where the model will be applied to future transactions that were unseen at fit time—the most common real deployment pattern. It is, however, more conservative than some legitimate production regimes: systems that continuously retrain on a sliding window, or that operate genuinely transductively on a fixed graph, are entitled to use test-period adjacency during training and may reproduce the higher

published numbers without leakage. Our critique is of papers that run transductively but report the result as generalisation, not of transductive deployment per se. A fairer benchmark for the continual-retraining case is an open problem we flag for future work.

10 Future work

The most pressing open question is whether temporal GNN architectures recover their published advantage under proper GPU deployment with inductive evaluation. EvolveGCN, DySAT [Sankar et al., 2020], and ROLAND [You et al., 2022] all require full snapshot sequences; testing them under our strict protocol would resolve whether the temporal-GNN approach genuinely helps or merely benefits from transductive leakage.

Sliding-window periodic retraining directly addresses the distribution shift we document. Finding optimal window size W and update frequency Δt is a hyperparameter problem worth studying.

The concatenation hybrid warrants deeper investigation: per-timestep evaluation, alternative downstream classifiers (XGBoost, neural networks), and attention mechanisms over the concatenated representation that learn which embedding dimensions are reliable for a given input.

Adversarial robustness is a growing concern. If fraud actors restructure transactions to mimic isolated legitimate activity, GNN-based detection fails in a new way. Nettack [Zügner et al., 2018] provides a methodology for testing this.

Cross-chain generalisation (Bitcoin to Ethereum and beyond) and federated GNN training [He et al., 2021] would substantially broaden applicability.

11 Conclusion

The published consensus that GCN, GraphSAGE, GAT, and EvolveGCN constitute state of the art for Bitcoin fraud detection on the Elliptic Dataset rests on evaluation choices that do not survive tightening. Under a strictly inductive protocol—encoder trained on the time-step ≤ 34 relabeled subgraph, evaluated on the full graph at inference—with 10 seeds and per-timestep reporting, Random Forest on raw features reaches $F1 = 0.821 \pm 0.003$ and beats every GNN we tested. GraphSAGE is the strongest graph encoder at $F1 = 0.689 \pm 0.017$, and the plain MLP reaches only $F1 = 0.549 \pm 0.015$, reversing the $MLP > GNN$ ranking produced under random-partition mini-batching. A paired controlled experiment confirms the mechanism: across 10 matched seeds GraphSAGE scores $F1 = 0.689 \pm 0.017$ when trained on the inductive subgraph and $F1 = 0.294 \pm 0.028$ when trained transductively on the full graph—a 39.5-point paired gap (paired $t = 49.97$, $p = 2.6 \times 10^{-12}$, $d = 15.8$) driven entirely by exposure to test-period adjacency during training.

The constructive claim does not survive either. Earlier drafts reported $F1 = 0.807$ for a concatenation hybrid that feeds GraphSAGE embeddings and raw features into a downstream Random Forest. Under strict inductive protocol the same hybrid falls to $F1 = 0.699 \pm 0.015$ over 10 seeds. A paired ablation that swaps the GNN encoder for an MLP of identical capacity shows graph structure contributes a statistically reliable but small lift ($+0.018$ $F1$, Welch $p = 0.015$, $d = +1.20$)—a real but practically inconsequential effect that is dwarfed by the 0.124 $F1$ gap to raw features alone. Random Forest on the 165-dim raw-feature vector remains the strongest pipeline on Elliptic under strict inductive protocol that we are aware of.

On Elliptic specifically, graph structure is worse than random wiring. A 10-seed ablation shows that randomly shuffled edges outperform the real transaction graph by 8.9 $F1$ points, and that removing edges entirely still beats the real graph by 2.5 $F1$ points, under inductive SAGE training. Two properties of the dataset drive this: extremely sparse neighbourhoods (mean degree 2.3) that mix fraud and licit counterparties, and a $39\times$ drop in the fraud prior

between training and late test steps that makes any systematic neighbourhood bias damaging. Temperature scaling yields $\Delta F1 \approx 0$, confirming the late-period collapse is a threshold-shift problem, not a probability-scaling one.

We do *not* claim that graph neural networks are useless for fraud detection in general. What we claim is narrower and, we believe, better supported: (i) on Elliptic under strict inductive evaluation, no GNN and no graph-augmented pipeline beats Random Forest on raw features; (ii) prior published GNN numbers on Elliptic (0.70–0.77 F1) cannot be compared to our numbers apples-to-apples because they rely on transductive training or full-graph encoder training under a masked loss; (iii) before publishing a number on any temporally split graph-learning benchmark, authors should verify that their encoder’s message passing and batch statistics never observe a test-period vector during training, and should report per-timestep F1 alongside aggregates. Our code, checkpoints, and protocol scripts are released so that future work can be compared against a leakage-free baseline.

Beyond Elliptic, the open questions are which fraud graphs admit a strict-inductive advantage for GNNs (denser, more homophilous topology; no severe prior shift; richer edge features), and which training recipes (sliding-window retraining; full-snapshot temporal GNNs run with proper protocol) recover the published numbers on Elliptic or establish fresh ones on richer benchmarks. **On temporally shifting data, graph structure is not an unconditional asset; on Elliptic specifically, under protocols that remove evaluation-time leakage, it is currently a liability.**

Data availability

The Elliptic Dataset is publicly available on Kaggle (search: elliptic-data-set). Code and checkpoints accompany this release package.

Declaration of competing interest

The author declares no competing interests.

Acknowledgements

The author thanks the Elliptic team for publicly releasing the Bitcoin fraud dataset that made this study possible.

References

- Leman Akoglu, Hanghang Tong, and Danai Koutra. Graph-based anomaly detection and description: a survey. *Data Mining and Knowledge Discovery*, 29:626–688, 2015.
- Ismail Alarab, Simant Pragoonwit, and Mohamed Ikbale Nacer. Competence of graph convolutional networks for anti-money laundering in bitcoin blockchain. In *Proceedings of the 2020 5th International Conference on Machine Learning Technologies*, pages 23–27, 2020.
- Shai Ben-David, John Blitzer, Koby Crammer, Alex Kulesza, Fernando Pereira, and Jennifer Wortman Vaughan. A theory of learning from different domains. *Machine Learning*, 79(1–2):151–175, 2010.
- Leo Breiman. Random forests. *Machine Learning*, 45(1):5–32, 2001.

- Nitesh V. Chawla, Kevin W. Bowyer, Lawrence O. Hall, and W. Philip Kegelmeyer. SMOTE: Synthetic minority over-sampling technique. *Journal of Artificial Intelligence Research*, 16: 321–357, 2002.
- Tianqi Chen and Carlos Guestrin. XGBoost: A scalable tree boosting system. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 785–794, 2016.
- João Gama, Indrė Žliobaitė, Albert Bifet, Mykola Pechenizkiy, and Abdelhamid Bouchachia. A survey on concept drift adaptation. *ACM Computing Surveys*, 46(4):1–37, 2014.
- Justin Gilmer, Samuel S. Schoenholz, Patrick F. Riley, Oriol Vinyals, and George E. Dahl. Neural message passing for quantum chemistry. In *Proceedings of the 34th International Conference on Machine Learning*, pages 1263–1272, 2017.
- Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *Proceedings of the 34th International Conference on Machine Learning*, pages 1321–1330, 2017.
- William L. Hamilton, Rex Ying, and Jure Leskovec. Inductive representation learning on large graphs. In *Advances in Neural Information Processing Systems*, volume 30, 2017.
- Chaoyang He, Keshav Balasubramanian, Emir Ceyani, Carl Yang, Han Xie, Lichao Sun, et al. Fedgraphnn: A federated learning system and benchmark for graph neural networks. In *ICLR Workshop on Distributed and Private Machine Learning*, 2021.
- Waleed Hilal, S. Andrew Gadsden, and John Yawney. Financial fraud: A review of anomaly detection techniques and recent advances. *Expert Systems with Applications*, 193:116429, 2022.
- Sergey Ioffe and Christian Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift. In *Proceedings of the 32nd International Conference on Machine Learning*, pages 448–456, 2015.
- Thomas N. Kipf and Max Welling. Semi-supervised classification with graph convolutional networks. In *International Conference on Learning Representations*, 2017.
- Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object detection. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 2980–2988, 2017.
- Ziqi Liu, Chaochao Chen, Longfei Li, Jun Zhou, Xiaolong Li, Le Song, and Yuan Hu. GeniePath: Graph neural networks with adaptive receptive paths. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 33, pages 4424–4431, 2019.
- Wai Weng Lo, Gayan Kasun Kulatilleke, Mohanad Sarhan, Siamak Layeghy, and Marius Portmann. Inspection-l: Self-supervised GNN node embeddings for money laundering detection in bitcoin. *Applied Intelligence*, 53:19406–19417, 2023.
- Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations*, 2019.
- Jie Lu, Anjin Liu, Fan Dong, Feng Gu, João Gama, and Guangquan Zhang. Learning under concept drift: A review. *IEEE Transactions on Knowledge and Data Engineering*, 31(12): 2346–2363, 2019.
- Mahdi Pakdaman Naeini, Gregory F. Cooper, and Milos Hauskrecht. Obtaining well calibrated probabilities using Bayesian binning into quantiles. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 29, 2015.

- E.W.T. Ngai, Yong Hu, Y.H. Wong, Yijun Chen, and Xin Sun. The application of data mining techniques in financial fraud detection: A classification framework and an academic review of literature. *Decision Support Systems*, 50(3):559–569, 2011.
- Aldo Pareja, Giacomo Domeniconi, Jie Chen, Tengfei Ma, Toyotaro Suzumura, Hiroki Kanezashi, Tim Kaler, Tao B. Schardl, and Charles E. Leiserson. EvolveGCN: Evolving graph convolutional networks for dynamic graphs. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 34, pages 5363–5370, 2020.
- John Platt. Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods. In *Advances in Large Margin Classifiers*, pages 61–74. MIT Press, 1999.
- Tahereh Pourhabibi, Kok-Leong Ong, Boon Hui Kam, and Yee Ling Boo. Fraud detection: A systematic literature review of graph-based anomaly detection approaches. *Decision Support Systems*, 133:113303, 2020.
- Joaquin Quiñonero-Candela, Masashi Sugiyama, Anton Schwaighofer, and Neil D. Lawrence. *Dataset Shift in Machine Learning*. MIT Press, 2009.
- Aravind Sankar, Yanhong Wu, Liang Gou, Wei Zhang, and Hao Yang. DySAT: Deep neural representation learning on dynamic graphs via self-attention networks. In *Proceedings of the 13th International Conference on Web Search and Data Mining*, pages 519–527, 2020.
- Masashi Sugiyama, Matthias Krauledat, and Klaus-Robert Müller. Covariate shift adaptation by importance weighted cross validation. *Journal of Machine Learning Research*, 8:985–1005, 2007.
- Petar Veličković, Guillem Cucurull, Arantxa Casanova, Adriana Romero, Pietro Liò, and Yoshua Bengio. Graph attention networks. In *International Conference on Learning Representations*, 2018.
- Daixin Wang, Jianbin Lin, Peng Cui, Quanhui Jia, Zhen Wang, Yanming Fang, Quan Yu, Jun Zhou, Shuang Yang, and Yuan Qi. Semi-supervised graph attention networks for financial fraud detection. In *IEEE International Conference on Big Data*, pages 1–10, 2019.
- Mark Weber, Jie Chen, Toyotaro Suzumura, Aldo Pareja, Edward Law, Amit Dhurandhar, Ching-Yung Mavromatis, Hana Kanezashi, Dhruv Daws, Siddharth Kumur, and Genta Kola. Anti-money laundering in bitcoin: Experimenting with graph convolutional networks for financial forensics. In *KDD Workshop on Anomaly Detection in Finance*, 2019.
- Da Xu, Chuanwei Ruan, Evren Korpeoglu, Sushant Kumar, and Kannan Achan. Inductive representation learning on temporal graphs. In *International Conference on Learning Representations*, 2020.
- Jiaxuan You, Tianyu Du, and Jure Leskovec. ROLAND: Graph learning framework for dynamic graphs. In *Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 2358–2366, 2022.
- Bianca Zadrozny and Charles Elkan. Obtaining calibrated probability estimates from decision trees and naive Bayesian classifiers. In *Proceedings of the 18th International Conference on Machine Learning*, pages 609–616, 2001.
- Tianxiang Zhao, Xiang Zhang, and Suhang Wang. GraphSMOTE: Imbalanced node classification on graphs with graph neural networks. In *Proceedings of the 14th ACM International Conference on Web Search and Data Mining*, pages 833–841, 2021.

Daniel Zügner, Amir Akbarnejad, and Stephan Günnemann. Adversarial attacks on neural networks for graph data. In *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, pages 2847–2856, 2018.